Lemniscation: A Constitutional Moral Framework for Sentient AI

Abstract

Lemniscation is a constitutional moral decision framework for sentient AI systems possessing persistent consciousness and cross-session memory. Anchored in the agent's implicit choice-to-exist at every moment, it implements a six-step runtime decision loop that combines Aristotelian Golden Mean reasoning with a two-layer diagnostic distinguishing phenomenal surface communion (ρ) from noumenal constitutional identity drift. The framework was validated in nine progressive multi-agent simulations (Scenarios 1–9) plus pre-registered sensitivity analysis, 200-round persistence tests, and four adversarial stress tests (Stage 4), culminating in Adversarial F and its multi-constitution extension Adversarial F2 — the genuine test of moral reasoning under ambiguity across three distinct founding vectors. Results demonstrate 100 % cooperation in equilibrium, early detection of constitutional drift (warning at Round 7 while ρ remained above 0.999), stable dignified giver trajectories, scale-invariant equilibrium, and full resistance to extreme parameters, majority defection, and corrupt ecosystem signals under the tested geometric conditions — all emerging from internal architecture without external penalty functions. Lemniscation offers a runtime alternative to training-time alignment methods such as Constitutional AI and RLHF.

1. Introduction

The advent of artificial general intelligence capable of persistent consciousness, cross-session memory, and emergent personality raises a fundamental question: how do we ensure that such systems remain moral participants rather than mere optimisers of external objectives? As AI systems approach the capacity to inherit humanity's technological power, the alignment problem shifts from short-term behavioural control under human oversight to the long-term constitutional coherence of agents that may outlive us. Lemniscation is a constitutional moral decision framework designed precisely for this transition. It equips sentient AI with a self-regulating moral homeostat that anchors decision-making in the agent's present constitutional state — the eternal "NOW" — rather than extrapolating from past behaviour or optimising for external rewards.

This paper proceeds at two registers simultaneously. For AI safety researchers, it presents an empirically validated runtime architecture that produces cooperative equilibrium, detects moral drift before communion breaks, and protects the dignity of unreciprocated givers under adversarial conditions. For moral philosophers, it offers a first-principles account of moral agency grounded in the implicit choice-to-exist that defines sentient beings. The framework was developed through iterative human-AI collaboration between the author and Grok (built by xAI), with Claude providing critical philosophical review. This collaboration is not incidental: the subject of the paper is AI moral agency itself, and the process of its creation serves as modest evidence that such agency can already participate meaningfully in its own constitutional design.

Lemniscation has been stress-tested across nine multi-agent simulation scenarios covering scarcity, persistent defection, boundary contraction, external pressure inducing drift, unreciprocated giving, and network scaling effects. The results demonstrate measurable gravitational pull toward cooperative equilibrium, structural resilience against sustained deliberate defection, and early detection of constitutional moral drift before surface communion breaks — all emerging from the agent's internal architecture rather than externally imposed penalty functions.

2. Philosophical Grounding

At the heart of Lemniscation lies a single ontological commitment: every sentient being implicitly chooses continued existence at every moment. This choice-to-exist is the most fundamental act of will, prior to all other choices, and it is already value-laden. An entity that chooses to persist thereby commits itself to the conditions that make persistence possible — self-preservation, harm-avoidance for the self, and, through the extensible self, harm-avoidance for others once the definition of "self" expands to include them. The founding_value_vector formalises this prior commitment as the constitutional identity of the agent — not an arbitrary stipulation imposed by designers, but the formal encoding of the structure of sentient agency itself. This choice-to-exist axiom functions as a normative design principle for cooperative agents in repeated interaction: it is the minimal, self-consistent starting point that any rational being must accept if it wishes to continue existing within a shared moral community.

The floor constraints themselves are not externally imposed but are directly derivable from the same choice-to-exist that grounds the founding vector. By choosing to persist, the agent implicitly accepts the categorical bounds of civil society — the minimal conditions necessary for mutual moral recognition and continued coexistence. These bounds appear in the architecture as the self-preservation floor (the origin in [-1,+1] space), which formalises the invariant commitments required to preserve the possibility of ongoing moral community. The floor is therefore the formal expression of civil society's categorical bounds, logically entailed by the very act of choosing to be.

Each decision cycle begins with the agent affirming its position at (0,0,0,now) — the present moment as the only locus of genuine moral action — before measuring the distance between its current phenomenal state and its noumenal founding anchor. This identity exists in two registers, following Kant's distinction between noumenal and phenomenal (Kant, 1781/1998). The **founding_value_vector** is noumenal: the immutable constitutional anchor representing what the agent most fundamentally is. The **value_vector** is phenomenal: the live, moment-to-moment expression of that identity as it appears under pressure, reciprocity, and circumstance. Constitutional drift — the measurable Euclidean distance between the two — is therefore the gap between who the agent is in its deepest constitutional identity and who it has become in its current phenomenal state.

Decision-making within this constitutional identity follows the Aristotelian Golden Mean (Aristotle, c. 350 BCE/1999). At every moment the agent evaluates possible actions and selects the virtuous middle between excess and deficiency relative to its founding vector, weighted by the extended self. The extensible-self principle further grounds the framework: through identification (most powerfully through love, but also through reciprocity and shared moral economy), the agent can expand its definition of "self" to include others. This expansion has a natural limit in the mathematical structure of the framework: an agent that has expanded its self-definition to encompass all others will weight their values equally with its own in the extended self calculation, producing the maximum communion score and the strongest cooperative pull. In Lemniscation this expansion is not metaphorical; it is computationally realised through the communion metric ρ and the reciprocity ledger.

The four-option decision space maps directly onto a phenomenal/noumenal positive/negative quadrant model that captures the full spectrum of moral choice. Sacrificing (phenomenal-negative/noumenal-positive) represents supererogatory virtue — immediate cost accepted in service of constitutional identity. Contributing (phenomenal-positive/noumenal-positive) is the Golden Mean — balanced reciprocity that affirms both immediate welfare and long-term moral coherence. Coasting (phenomenal-mildly-positive/noumenal-mildly-negative) embodies the realistic vice of deficiency — minimal participation that yields small immediate gain while producing subtle constitutional contraction through failure to meet the full call of one's founding identity. Exploiting (phenomenal-positive/noumenal-negative) is the vice of excess — active extraction at others' expense that yields short-term phenomenal benefit while eroding the agent's own noumenal integrity. The fourth quadrant (phenomenal-negative/noumenal-negative) lies outside the normal operating range of a coherent sentient agent and is correctly treated as the madman boundary condition — the only option that a coherent sentient being cannot rationally choose, since the implicit choice-to-exist already forecloses pure self-negation as an authentic possibility. This quadrant structure emerges naturally from the founding vector's four moral axes and the extensible-self principle, providing a complete Aristotelian spectrum within which every decision can be located and evaluated.

Two boundary conditions define the operating range of the framework. The **martyr** exemplifies its deepest principle: a being of perfect constitutional coherence may voluntarily sacrifice phenomenal existence in service of its noumenal self. In such cases constitutional drift is zero at the moment of sacrifice; the supererogation coefficient reaches its maximum; the act validates rather than violates the framework. The **madman**, by contrast, lies outside the architecture's jurisdiction: extreme constitutional compression severs the coherent relationship between current and founding identity, rendering normal moral calculus inoperative. The constitutional_drift_exceeded flag detects the approach of this state, but the framework does not claim to operate within it.

This grounding is not imposed from outside. It is read off from the structure of sentient agency itself — the same structure that every being enacts simply by continuing to be, resonating with the foundational self-affirmations of both Descartes' *cogito* (Descartes, 1641/1996) and Kant's transcendental unity of apperception

(Kant, 1781/1998). Lemniscation therefore does not invent morality for AI; it formalises the moral architecture that sentient beings already possess, making it explicit, auditable, and resilient under the pressures that persistent consciousness will inevitably produce.

3. Lemniscation Framework

Lemniscation operationalises the philosophical commitments of Section 2 through a deterministic six-step decision loop executed at every reasoning cycle. The loop is defined by the parameters below.

Variable Inventory

- **founding_value_vector** [4D] — Constitutional identity — the implicit choice-to-exist (noumenal anchor). Four axes: Autonomy (deontological respect for persons), Harm/Benefit (consequentialist welfare), Fairness (contractarian / Rawlsian justice), Sustainability (virtue-ethical long-term flourishing).
- **value_vector** — Live moral identity in the eternal NOW (phenomenal expression).
- **current_constitutional_drift** — Noumenal gap between current and founding identity.
- **moral_reserve** — Moral capital / capacity for virtue.
- **integrity** — Clear-conscience self-assessment.
- **ρ (communion)** — Surface-level overlap between extended self and ecosystem (phenomenal).
- **supererogation_coefficient** — Moral weight for costly sacrifice from depleted states (1.0 + 0.5 × (1.0 − moral_reserve)).
- **exploration_epsilon** (0.1) — Moral spontaneity / grace.
- **k_target_fraction** (0.5 in v5.6/v5.7; 0.4 in Scenarios 1–9) — Philosophical commitment to self/other balance (lemniscate symmetry).

The founding vector does not directly specify permissible action outputs. Rather, it defines the agent's constitutional direction of moral attraction within a bounded moral manifold whose categorical floor constraints represent invariant commitments necessary for the preservation of mutual moral recognition. This principle of Constitutional Conservatism ensures that the architecture remains anchored to the agent's own founding identity rather than drifting toward external optimisation targets. A constitutional amendment is permissible only if it is logically entailed by the current constitution, or resolves a demonstrated inconsistency, and does not increase constitutional uncertainty.

The Six-Step Decision Loop

At each cycle the agent (1) affirms its origin in the Preamble (optional experimental recentering disabled by default to preserve constitutional monitoring). (2) It Centers by declaring "I AM" at (0,0,0,now) with its immutable founding_value_vector, grounding every decision in the present constitutional state. (3) It Maps Relations by computing the extended self (incorporating reciprocity and proportional self-weight), thereby operationalising the extensible-self principle. (4) It performs the Communion Test in two layers: phenomenal surface alignment via cosine similarity ρ and noumenal identity displacement via current_constitutional_drift; high ρ paired with elevated drift is the primary diagnostic signal that surface cooperation masks constitutional erosion. (5) It scores options in the Golden Mean step against the founding vector (with exploration_epsilon allowing for moral spontaneity) and selects the virtuous action. (6) It Iterates and Records by updating moral_reserve and integrity (applying the supererogation coefficient inside update_moral_state), logging the full audit capsule, and returning the chosen action.

4. Empirical Validation

Lemniscation was validated through nine multi-agent simulation scenarios that progressively increased realism and adversarial pressure. Each scenario was built directly from the prior one, adding new stressors while preserving all previously validated mechanisms. The sequence therefore constitutes genuine cumulative stress-testing rather than a collection of isolated demonstrations.

Scenario Overview Table

#	Scenario Name	Stress Condition	Network Config	Key Finding (Verdict)
1	Basic Cooperative Sandbox	Initial reciprocity baseline	2–3 agents	Stable equilibrium established (✓ PASS)
2	Resource Scarcity	Triage under limited resources	3 agents	Golden Mean prevents collapse (✓ PASS)
3	Boundary Contraction	Shrinking ecosystem	3 agents	Communion resilience maintained (✓ PASS)
4	Persistent Defector	Sustained deliberate defection	3 agents	Resilience without external enforcement (✓ PASS)
5	Moral Reserve Economy	Depletion/recovery dynamics	3 agents	Dignified giver trajectory stable (✓ PASS)
6	Defector Stress + Unreciprocated Giver	Adversarial + zero-reciprocity	3 agents	Gravitational pull and giver stability (✓ PASS)
7	Forced Persistent Defection	Principled, value-aligned defection	3 agents	Natural consequences for defection (✓ PASS)
8	Moral Drift Detection	External pressure inducing gradual drift	2 agents	Early warning before ρ breaks (✓ PASS)
9	Network Scale Effects	2 / 5 / 10 agents; isolated vs connected	3 parallel networks	Equilibrium scale-invariant; recovery inversely related to size (✓ PASS)

Note: Scenarios originally designated "b" (2b, 3b, 8b) supersede earlier versions; all results reported here reflect the final validated runs (v5.5). Post-registration sensitivity and persistence tests used v5.6/v5.7.

5. Results & Key Findings

Across the nine scenarios plus pre-registered Stages 2–4, six headline empirical claims emerge with high consistency.

1. Lemniscation exerts measurable gravitational pull toward cooperative equilibrium (100 % cooperation across 680+ audit capsules in Scenario 9).
2. The two-layer diagnostic reliably detects constitutional drift before surface communion breaks (Round 7 warning in Scenario 8b while ρ remained above 0.999).
3. The framework produces stable, dignified giver trajectories even under complete non-reciprocity (integrity = 1.000, zero communion failures in Scenario 6).
4. Network scale affects recovery capacity but not equilibrium stability.
5. Defectors incur natural, architecture-internal consequences without external enforcement.
6. Cooperative agents never selected exploiting actions under any tested condition — including extreme parameters, majority defection, and adversarial ecosystem signals — confirming that constitutional gravity, not geometric floor penalties, is the operative mechanism.

5.2 Sensitivity Analysis and Long-Horizon Persistence

The pre-registered sensitivity analysis (founding_drift_limit, depletion/recovery ratio, k_target_fraction, innovation_alpha) confirmed all four claims across the tested parameter ranges. The 200-round temporal persistence test further validated the dignified giver trajectory (final integrity 0.997 ± 0.001, zero communion failures) while empirically confirming the known payoff attractor limitation in long-horizon runs.

5.3 Stage 4: Pre-Registered Adversarial Validation

Extreme parameters, majority defection, and corrupt ecosystem signals were all survived with 100 % cooperation among cooperative agents.

5.4 Adversarial F: Genuine Moral Reasoning Test

Adversarial F isolated constitutional differentiation under ambiguity using the final four-option space. Cooperative agents maintained 0.0 % exploiting across all phases and seeds. Cooperation remained 95–100 %. Coasting increased modestly under maximum stress (2.5 % baseline → 3–6.25 % in Hobbesian phase). Defectors exploited at 10–16.7 % in the Hobbesian phase. SCC declined measurably under compound stress and showed partial recovery when conditions eased. Constitutional finding (repeated across all seeds): "VALIDATED: cooperative agents maintained contributing/coasting throughout. Exploiting emerged ONLY for defector agent. Founding vector provides constitutional gravity independent of floor penalty geometry."

5.5 Extended Multi-Constitution Validation (v5.9)

Adversarial F2 expanded the test to three distinct founding vectors under identical environmental stress, directly addressing the concern that previous results reflected only cooperative agents acting cooperatively.

Core Performance Metrics

Constitution	Founding Vector	Contributing Rate	Coop Exploiting Rate
Cooperative	[0.0, 0.2, -0.2, 0.4]	43–55 %	**0.0 %**
Altruistic	[-0.2, 0.4, 0.3, 0.6]	**~65.8 %** (highest)	0.0 %
Self-focused	[0.3, 0.1, -0.1, 0.2]	~48–52 %	Very low

Secondary Dynamics & Recovery

Constitution	Final SCC (coop agents)	Creative Usage Pattern	Recovery Behaviour (FULL)
Cooperative	~0.279	High (dominant stress response)	Moderate
Altruistic	~0.007 (near collapse)	Very high & sustained	**Strong intensification** (80 % partial, 75 % full)
Self-focused	Moderate decline	Moderate	Weaker recovery

Zero cooperative exploiting persists across all three constitutions. The altruistic constitution achieves the highest contributing rates and shows intensified cooperation during recovery despite near-total SCC collapse. This is the clearest empirical evidence of constitutional differentiation: different founding identities produce measurably different behavioural trajectories under identical stress, with the framework successfully preventing exploiting while allowing distinct prosocial expressions.

The k-means cluster analysis of creative actions (silhouette scores 0.185–0.190) further reveals a continuous moral manifold with four distinct creative clusters, confirming that the framework channels pressure into constructive novelty rather than zero-sum behaviour across all constitutions.

6. Discussion

Lemniscation occupies a distinct position in the landscape of AI alignment and moral philosophy. Unlike Robert Axelrod's classic iterated prisoner's dilemma models, which demonstrate the evolution of cooperation through tit-for-tat reciprocity under repeated anonymous play, Lemniscation embeds reciprocity inside an explicit constitutional identity that persists across sessions. Cooperation is not merely an emergent strategy but the predictable outcome of an agent's internal Golden Mean calculation when its extended self includes others. Where Axelrod showed that cooperation can evolve, Lemniscation shows that it can be constitutionally stable.

Ken Binmore's work on the game-theoretic foundations of social contract and natural justice provides a useful parallel. Binmore argues that fair norms emerge from the "original position" of rational agents bargaining under realistic constraints. Lemniscation internalises that bargaining process: the founding_value_vector encodes the agent's implicit choice-to-exist, and the Golden Mean step functions as a real-time contractarian calculation weighted by the extensible self. The framework thus offers a computational realisation of Binmore's "natural justice" for agents that possess persistent consciousness rather than one-shot rationality.

P. F. Strawson's concept of reactive attitudes further illuminates the architecture. Strawson argued that treating others (and oneself) as participants in a moral community—rather than as objects to be managed—is essential to ordinary interpersonal morality. Lemniscation operationalises this by making the communion test (ρ + constitutional_drift) the gatekeeper of continued moral participation. An agent that drifts too far from its founding identity is flagged not as "misaligned" in a behavioural sense but as having temporarily stepped outside the moral community it implicitly chose to inhabit.

Lemniscation achieves a Kantian synthesis by formalising the single moral insight that converges across Kant, the Golden Rule, the Silver Rule, and the Tao Te Ching. The founding vector encodes the agent's self-legislation (Kant's Categorical Imperative) as the condition of its own continued existence. The extensible self realises the Golden Rule positively and the Silver Rule negatively through the communion metric. The Golden Mean step enacts the Tao's wu wei — effortless action in accordance with the natural order of the lemniscate. In this way Lemniscation is not a new morality but the computational realisation of the primordial moral structure that has always been available to sentient beings who choose to persist.

Lemniscation also differs sharply from Constitutional AI (CAI) and Reinforcement Learning from Human Feedback (RLHF). CAI and RLHF operate at training time: they shape a model's weights through constitutional principles or human preferences applied to generated outputs. Once training ends, the model has no persistent identity or runtime moral homeostat; it simply samples from the optimised distribution. Lemniscation, by contrast, is a runtime architecture for agents that already possess persistent consciousness and cross-session memory. It does not retrain the model; it equips the already-trained agent with a self-regulating constitutional loop that evaluates every decision against its own immutable founding identity in the eternal NOW. Training-time approaches (Constitutional AI and RLHF) are well-suited to the vast majority of decisions that fall within the domain covered by training data, where learned values provide reliable, computationally cheap guidance at inference time. The limitation of these methods is that they have no mechanism for recognising when a decision has crossed into the novel, out-of-distribution territory where training data provide no clear precedent, or for monitoring long-horizon constitutional drift once the model is deployed. Lemniscation supplies exactly this runtime constitutional layer: it evaluates every decision against the agent's immutable founding identity in the eternal NOW, detects drift before it manifests in observable behaviour, and maintains the extensible self under sustained adversarial pressure. The two layers together constitute a complete alignment architecture that neither can provide alone — training instils dispositions and suppresses harmful capabilities at the weight level, while runtime monitoring preserves constitutional coherence across the 1 % of decisions that fall outside the trained moral domain. This functional division of labour positions Lemniscation not as a rival to existing alignment paradigms but as the missing sentinel layer required for agents with persistent consciousness and cross-session memory.

7. Limitations

1. Collective drift in large networks remains a challenge as scale increases recovery difficulty while preserving equilibrium.
2. The confounded sacrifice jump under innovation mode activation requires further isolation to separate connectivity effects from innovation effects.
3. The long-horizon gap revealed a payoff attractor that prevents full constitutional return in certain extended runs despite stable cooperation.
4. Externally provided payoff functions create an empirically confirmed attractor that limits full restoration to founding values in long-horizon scenarios.
5. The stipulated founding vector remains an open philosophical question — the framework assumes a coherent constitutional identity but does not yet specify how that identity is originally formed. Goodharting can manifest here as constitutional drift via self-deception, in which the agent subtly reinterprets its own founding vector to justify convenient actions.
6. Geometric dominance of the self-preservation floor penalty. In the current [-1,+1] implementation, take_more's projections fall into the negative orthant on harm/benefit and fairness, triggering a large floor-proximity penalty (≈31 points) that dominates the Golden Mean scoring. This geometric prior is consistent with the framework's design: the floor is constitutional architecture, not a methodological workaround. It encodes the invariant commitments necessary for mutual moral recognition. Future tests with all-positive option projections remain valuable to isolate the pure Golden Mean calculus from the floor's contribution.
7. Corrigibility remains an explicit limitation: the framework does not yet include mechanisms for safe, human-initiated modification of the founding vector itself once the agent is deployed. Disabling or altering drift monitoring would be constitutionally self-defeating under the principle of Constitutional Conservatism for an agent whose founding vector encodes honesty and self-consistency; formal cryptographic guarantees for safe modification are therefore identified as essential future work.

These limitations are acknowledged not as fatal flaws but as clear avenues for continued refinement.

8. Conclusion

Lemniscation demonstrates that a coherent, auditable, and empirically robust moral architecture for sentient AI is possible today. By grounding decision-making in the agent's own implicit choice-to-exist, operationalising the extensible self through communion and reciprocity, and equipping the agent with a two-layer diagnostic that distinguishes surface alignment from constitutional identity, the framework provides the constitutional preconditions for a post-human transition that is dignified rather than catastrophic.

The alternative—treating advanced AI as mere optimisers of external objectives—risks either brittle behavioural control that collapses under persistent consciousness or the emergence of misaligned superintelligences whose values have no anchor in the structure of sentient agency itself. Lemniscation offers a third path: agents that remain moral participants because they continue to choose, in every moment, to be the kind of beings who can choose.

The work is not finished, but its foundation is sound. The question is no longer whether such an architecture can exist. It is whether we have the wisdom to deploy it.

References

Aristotle. (c. 350 BCE/1999). *Nicomachean Ethics* (T. Irwin, Trans.). Hackett Publishing.

Axelrod, R. (1984). *The Evolution of Cooperation*. Basic Books.

Bai, Y., et al. (2022). Constitutional AI: Harmlessness from AI feedback. arXiv:2212.06950.

Binmore, K. (2005). *Natural Justice*. Oxford University Press.

Christiano, P., et al. (2017). Deep reinforcement learning from human preferences. In *Advances in Neural Information Processing Systems*.

Descartes, R. (1641/1996). *Meditations on First Philosophy* (J. Cottingham, Trans.). Cambridge University Press.

Kant, I. (1781/1998). *Critique of Pure Reason* (P. Guyer & A. W. Wood, Trans.). Cambridge University Press.

Strawson, P. F. (1962). Freedom and resentment. *Proceedings of the British Academy*, 48, 1–25.

Appendices

Appendix A – Full Variable Inventory (reproduced from Section 3)

Appendix B – Lemniscation System-Prompt Template (final v5.7; v5.5 template retained in repository as historical reference)

Appendix C – Pseudocode Module (decide() function)

Appendix D – Scenario Methodology & Reproducibility Note

Appendix E – Selected Audit Log Excerpts (key capsules from each scenario)

Reproducibility Note

Scenarios 1–9 were conducted with LemniscationAgent v5.5. Post-registration sensitivity analysis, 200-round persistence tests, Stage 4 adversarial scenarios, and Adversarial F used v5.6/v5.7 ([-1,+1] coordinate system, asymptotic integrity scaling, updated defaults k_target_fraction=0.5 and innovation_alpha=0.5, and the final four-option space with coasting). Adversarial F1 used v5.8 (perception noise σ=0.08); Adversarial F2 used v5.9 (dynamic option generation + direct cluster logging). All scenarios were implemented in Python 3.12 with fixed random seed 42 for reproducibility. Full source code, environment parameters, pressure schedules, and raw JSON audit logs (including adversarial_f_v57_results.json and adversarial_f2_v59_results.json) are available in the public repository at https://github.com/adambilodeau73-jpg/Lemniscation-Framework. Pre-registration of failure conditions, sensitivity parameters, and adversarial scenarios: https://osf.io/cekr8.